M5 Ultra · 80 GPU cores · 256 GiB

Qwen3.8-27B NVFP4 · MLX 0.32.3 · AC power · 19-token prompt / 32-token output

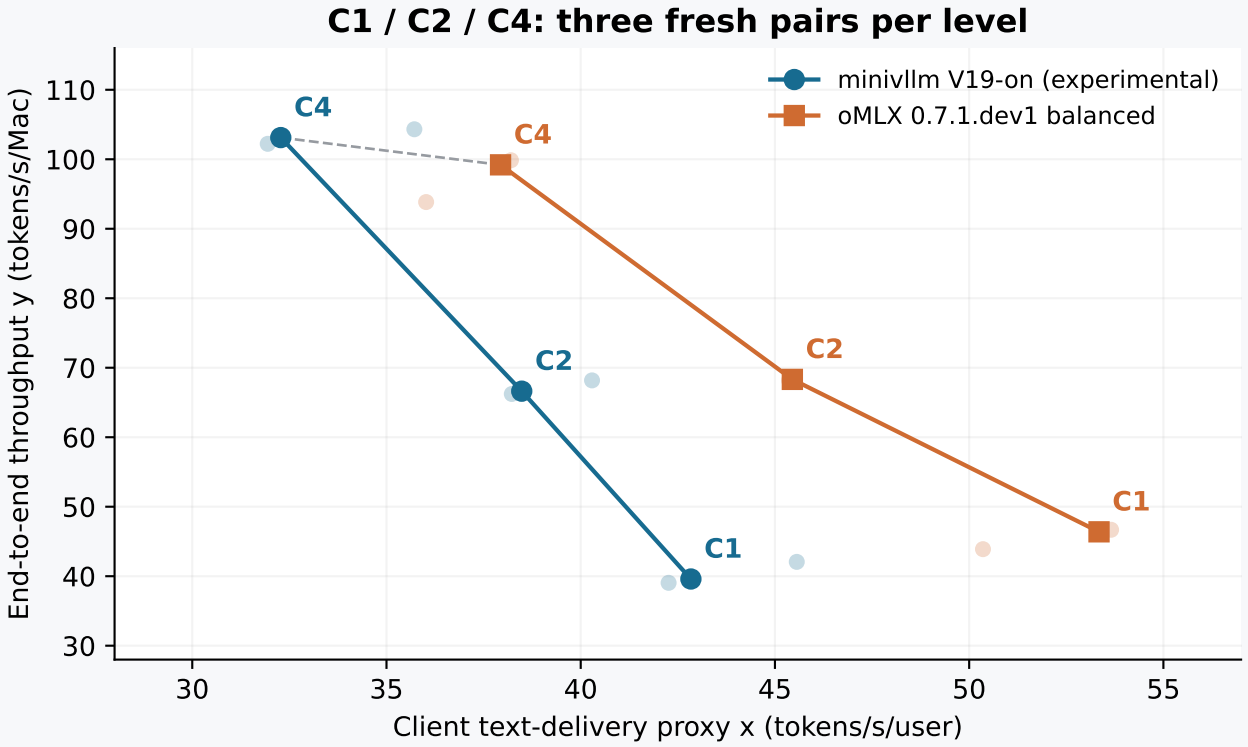

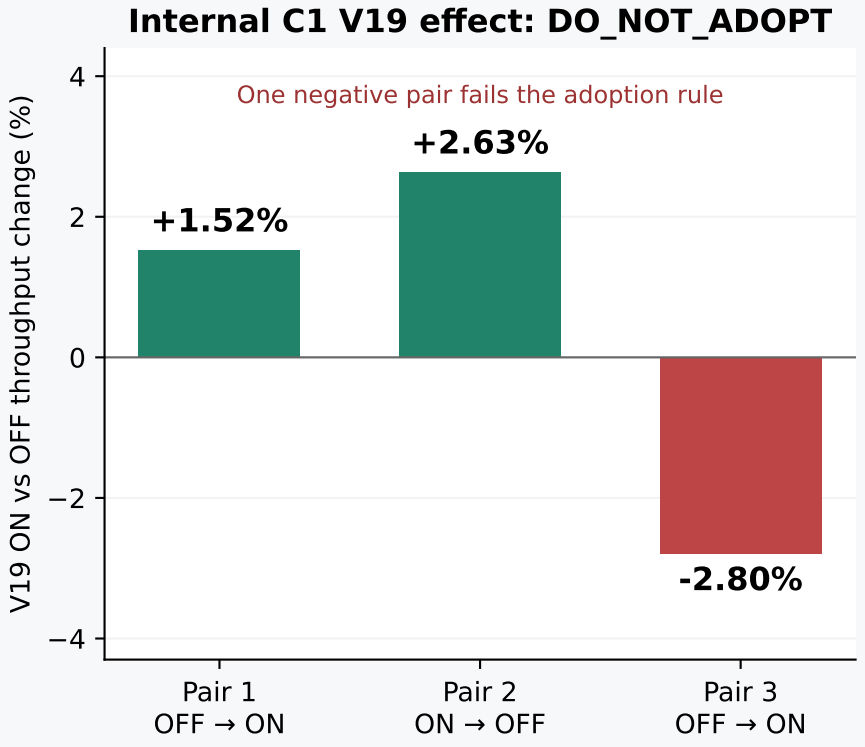

Left: median of three process medians; faint dots retain every process. Dashed line: observed Pareto frontier.
Right: a separate six-process OFF/ON experiment; aggregate y +2.46%, but pair 3 is negative.
x is a client delivery proxy, not GPU token rate. No M4 data pooled. Lithos has no supported M5 Ultra profile.